

# Databases for Data Science

Lecture 12 · 2022-10-12

# Today

- Discussion of other databases
  - SQLite
- PostgreSQL support for JSON
- Miscellaneous topics
- Extra practice

# All sorts of SQLs

Flavor	Developer	FOSS
<b>PostgreSQL</b>	UC Berkeley → Open-source contributors	✓
MySQL	Oracle	✗
MariaDB	MySQL fork → Open-source contributors	✓
SQL Server	Microsoft	✗
IBM Db2	IBM	✗
<b>SQLite</b>	Open-source contributors	✓

# SQLite

PostgreSQL and most other database engines have a standalone **server process**.

This runs separately from the clients, like the `psql` console and `psycopg2`, that connect to it.

SQLite, instead, is a **library** that works in the same process as your own code.

*some properties:*

	PostgreSQL	SQLite
Database engine runs in...	A standalone server process	The same process as your code
Storage	Multiple files on-disk	One <code>.db</code> file (or in-memory)
When you run a query...	Your client sends an HTTP request to the server	Library code is called directly
Write-concurrency	Isolated transactions	One writer at a time

# SQLite

Why use SQLite?

# SQLite

Why use SQLite?

- Easier to setup and use
- Runs in restricted environments (phones, embedded)
- Instances are isolated by default

So, why use a client/server system like Postgres?

# SQLite

Why use SQLite?

- Easier to setup and use
- Runs in restricted environments (phones, embedded)
- Instances are isolated by default

So, why use a client/server system like Postgres?

- Works with remote clients
- Better concurrency features
- Handles practically unlimited amounts of data

*further reading:* <https://www.sqlite.org/whentouse.html>

# SQLite: Practice

On CS1:

```
// copy the `orders.csv` file to your working directory
cp /usr/share/databases/orders.csv .

// start SQLite with a new database file
sqlite3 lec12.db
```

Next, in the SQLite console:

```
.import orders.csv raw_data

SELECT * FROM raw_data;
```

**Exercise:** In SQLite, split this data into `orders` and `customers` tables. Each order should have a foreign key to a customer.



# PostgreSQL JSON features

PostgreSQL added JSON support in v9.2 (2012).

- (SQLite followed suit in v3.38.0 - earlier this year!)

It's still a relational database, so how does this work?

- Tables can have `json`- and `jsonb`-typed columns
  - `json` is stored as plaintext, while `jsonb` is more efficient
- We have new operators and functions for working with JSON data in SQL.

# JSON in PSQL

```
CREATE TABLE products (  
    product_id SERIAL PRIMARY KEY,  
    title TEXT,  
    info JSONB  
);  
INSERT INTO products (title, info)  
VALUES  
    (  
        'Chalk', -- We use single quotes to start/end the JSON and double quotes inside it  
        '{ "manufacturer": "Crayola", "sizes": [12,24,48], "unit_price": 0.20 }'  
    ),  
    (  
        'USB Cable',  
        '{  
            "manufacturer": "Anker",  
            "skus": [  
                {"length": 1.5, "color": "black"},  
                {"length": 6.0, "color": "white"}  
            ]  
        }'  
    );
```

**Exercise:** On your own computer, create a `products` table and insert some records with JSON.

# JSON in PSQL : Queries

The arrow operator allows us to extract nested properties from the JSON.

```
SELECT
  title,
  info -> 'manufacturer' AS manufacturer,
  info -> 'skus' AS skus,
  info -> 'skus' -> 0 AS sku_0,
FROM products;
```

Properties that are missing will be mapped to `NULL` values.

## Exercise

Insert a product with a deeply nested property, then write a query to select that property.

## Exercise

Insert some products with `color` and `price` fields, e.g. `'{"color": "blue", "price": 12.99 }'`

Take the average price grouped by color.

# JSON in PSQL : Operators

Two operators:

- `->` returns a JSON data structure
- `->>` converts its output to a text string

```
WITH step1 AS (  
    SELECT  
        title,  
        info -> 'manufacturer' AS mfg_json,  
        info ->> 'manufacturer' AS mfg_text  
    FROM products  
)  
SELECT  
    title,  
    mfg_json,  
    pg_typeof(mfg_json),  
    mfg_text,  
    pg_typeof(mfg_text)  
FROM step1;
```

# JSON in PSQL : Arrays

We have multiple options for dealing with arrays.

Using specific indices:

```
SELECT
  title,
  info -> 'manufacturer' AS manufacturer,

  -- Select array elements by index
  info -> 'skus' -> 0 AS sku_0,
  info -> 'skus' -> 1 AS sku_1,
FROM products;
```

Using functions to transform the array:

```
SELECT
  title,
  info -> 'manufacturer' AS manufacturer,

  -- Unroll the array into separate rows
  json_array_elements(info -> 'skus') as sku,

  -- We can then process the individual array elements
  json_array_elements(info -> 'skus') -> 'color' as color
FROM products;
```

## Exercise

Create a new SKU table with columns `sku_id SERIAL PRIMARY KEY`, `color TEXT`, `info JSONB`.  
Write a query that selects SKUs from `products` and inserts them into the new table.

# PostgreSQL vs MongoDB

Postgres is competitive with Mongo in performance.

So, why use MongoDB?

# PostgreSQL vs MongoDB

Postgres is competitive with Mongo in performance.

So, why use MongoDB?

- Preference for its query language / workflows
- Working with highly unstructured data
- Better support for *sharding*
  - Distributing data across a wide set of computers

*further reading:* <https://www.mongodb.com/compare/mongodb-postgresql>

# Other NoSQL options

In general, more diversity than in the SQL world.

Flavor	Developer
MongoDB	Document store
CouchDB	Document store
HBase	Column-oriented
Neo4j	Graph database
DynamoDB	Key/value store (cloud)
Redis	Key/value store (in-memory)



## Note: The CAP Theorem

“ Consistency, availability, and partition tolerance - pick two. ”

- Consistency: All replicas appear to have the same state.
- Availability: Operations are always possible.
- Partition Tolerance: The system remains consistent if replicas are disconnected.

This is true of distributed systems in general.

# Open Q&A / Exercises

thank you!